

Research Statement

Mukesh Tiwari

June 30, 2020

My research focuses on formal verification of software programs used in democratic process, e.g. elections. The rationale for verifying these software programs is that various critical decisions are taken by government based on the output produced by these software programs. However, if the software program involved in the decision making process is bug ridden, then it may produce a wrong output which could lead to a poor decision making. As a result, these kind of poor decision making could hamper the trust of members of general public in system despite the fact that there were no such intentions by government. Worse yet, now a days there are more and more automation of decision making based on machine learning algorithm, artificial intelligence, etc. Therefore, it is more imperative than ever that we formally verify these software programs and develop software tools to automate the formal verification process make the decision making more trustworthy.

1 PhD Work

During my PhD, my research was focussed on verifying electronic voting, specifically vote counting schemes, and brining three important ingredients, correctness, privacy, and verifiability, of paper ballot election to electronic setting (electronic voting). In a paper ballot election, correctness and verifiability is assured by public scrutineers, appointed by participating candidates and political parties, observing the counting process conducted by election commission officials. Moreover, the privacy in paper ballot election comes for free because of the secret ballot mechanism, introduced by Australia in 1855, where the voters choices are anonymous. However, none of these desirable properties can be achieved in electronic voting because the software programs, used during the various stages of elections, works in very opaque way [1, 2]. This opacity can cause a harm of various level, but it would be devastating if the vote-counting software program produces a wrong winner, which is not intended by the voters, because of a software bug. In this setting, it is very difficult to establish correctness and verifiability because there is no (direct) involvement of human in the counting process (other than pressing some buttons to run the software program). In most cases, the vote-counting software programs lack the quality measures which could lead to various problems including, but not limited to, producing incorrect result, ballot identification, etc. More importantly, these software programs are treated as commercial in confidence and are not allowed to be inspected by members of general public [3]. As a consequence, the result produced by these software programs can not be substantiated, which is very detrimental for any democratic society. Because of these reasons, the early adopter of electronic voting system, Germany and The Netherlands, were the early abandoner [4].

In my thesis, I addressed the concern of correctness of vote counting software by implementing and proving the correctness of the implementation (Schulze Method) in Coq theorem prover. I also proved that my implementation of Schulze Method follows the various meta properties, e.g. Condorcet winner, reversal symmetry. I addressed the privacy concern by using homomorphic encryption to count the (encrypted) ballot without decrypting any individual ballot and verifiability concern by generating an independently checkable scrutiny sheet (certificate) augmented with zero-knowledge-proofs for various claims made during the counting [5, 6].

2 Current Work

Currently I am working as a Post Doctoral researcher with Toby Murray at the university of Melbourne, and I work on security concurrent separation logic for formally reasoning the information flow in concurrent programs. My current project focuses on a verified email server, inspired by Hagrid key-server, which leaks no sensitive information to attackers. The basic idea behind this project is that given all the APIs of an email server, an attacker would not be able to get any information which is classified sensitive (high). We have already presented our work in *VerifyThis 2020* challenge.

3 Future Work

Throughout my past and current research, I have extensively used the formal verification techniques to verify the software programs, and I would like to continue to do so in the future as well. In future, I would like to:

- focus on verifying machine learning algorithm. Now a days, machine learning algorithms are deployed in various domains and decisions are taken based on the output of these programs. In general, the most popular language to implement these algorithms are Python which has no concept of type safety. As a consequence, a programmer writing these implementation can make a subtle error, which could be catastrophic. In contrast, a type safe language, e.g. Haskell, would catch many of these errors at the compile time. However, none of these languages would give 100% guarantee that the implementation is correct. The only way we can get 100% guarantee by implementing these algorithm in a theorem prover, Agda, Coq, Lean, and proving that the implementation follows the specification.
- focus on pushing privacy preserving machine learning algorithm in public domain. Currently, all the deployed machine learning algorithms are trained on plaintext data. However, it poses a significant challenge if the training data contains sensitive information. A privacy preserving machine learning would have all the benefits of machine learning without any concern about leaking sensitive data to anyone. Right now, the field is in very nascent phase, but given that there is a lot of progress in homomorphic encryption since 2009, it would be an excellent idea to explore privacy preserving machine learning. The other possible direction could be multi-party computation, but these concepts are very difficult for machine learning researcher to understand. Developing a domain specific language which abstracts the underlying details

would help machine learning researchers to push the boundaries of privacy preserving machine learning.

- focus on designing probabilistic programming language, either embedded in a theorem prover, or a separate language, for cryptographic implementation with main focus on correctness. Cryptographic algorithms play a critical role in any application focussed on privacy. If any part of implementation is not done right, it could lead to severe consequences (possible identification of activist, voter). The long term (and ambitious) goal would be developing a verified compiler which would generate assembly code (having various properties such as memory safe, constant time). There are already some work in this area, e.g. Jasmin and CakeML; however, Jasmin's goal is to generate formally verified efficient assembly code and CakeML is a general purpose language. None of these are focussed on probabilistic programming which is a necessary component to model any cryptographic algorithm.
- focus on designing electronic voting protocol and proving the meta properties. Many countries, e.g. Switzerland, often conduct referendums via internet-voting on various issues that matter to their democracy, and it encourages their citizen to participate actively in decision making. It would be interesting to design and implement electronic voting protocols for low-coercion situations, which would encourage the voters to actively participate in various decision making process without any fear of vote selling.
- focus on information flow security in software design (continue my collaboration with Toby Murray). Proving formally the information flow between various components of a software program can leads to security focussed design. In certain cases, where security is paramount, it helps in understanding the potential information leak, which can further be used for assessing the repercussions.

Finally, My long-term aim is to make formal verification accessible and ubiquitous in software development, specifically deployed in public domain. I believe that my expertise in *Theorem Proving, Election Security, and Cryptography* gives me an unique perspective to solve challenging problem, which matters to many democracies.

References

- [1] Scott Wolchok, Eric Wustrow, J. Alex Halderman, Hari K. Prasad, Arun Kankipati, Sai Krishna Sakhamuri, Vasavya Yagati, and Rop Gonggrijp. Security analysis of india's electronic voting machines. In *Proceedings of the 17th ACM Conference on Computer and Communications Security, CCS '10*, pages 1–14, New York, NY, USA, 2010. ACM.
- [2] J. Alex Halderman and Vanessa Teague. The new south wales ivote system: Security failures and verification flaws in a live online election. In Rolf Haenni, Reto E. Koenig, and Douglas Wikström, editors, *E-Voting and Identity*, pages 35–53, Cham, 2015. Springer International Publishing.
- [3] Australian Electoral Commission. Letter to Mr Michael Cordover, LSS4883 Outcome of Internal Review of the Decision to Refuse your FOI Request no. LS4849, 2013. available via [http](http://):

[//www.aec.gov.au/information-access/foi/2014/files/ls4912-1.pdf](http://www.aec.gov.au/information-access/foi/2014/files/ls4912-1.pdf), retrieved October 23, 2019.

- [4] Bart Jacobs and Wolter Pieters. *Electronic Voting in the Netherlands: From Early Adoption to Early Abolishment*, pages 121–144. Springer Berlin Heidelberg, Berlin, Heidelberg, 2009.
- [5] Thomas Haines, Dirk Pattinson, and Mukesh Tiwari. Verifiable homomorphic tallying for the schulze vote counting scheme. In Supratik Chakraborty and Jorge A. Navas, editors, *Verified Software. Theories, Tools, and Experiments*, pages 36–53, Cham, 2020. Springer International Publishing.
- [6] Thomas Haines, Rajeev Goré, and Mukesh Tiwari. Verified verifiers for verifying elections. In *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security, CCS '19*, page 685–702, New York, NY, USA, 2019. Association for Computing Machinery.